DEEPSTACK The Buildout · Part 05 · Memo

AI LABS

OpenAI and Anthropic Are Not the Same Business

OpenAI is cutting prices while Anthropic still monetizes enterprise coding. The difference matters because hyperscalers, lenders and investors have priced the pair as one AI bet.

By DeepStack · September 11, 2026 · AI Labs · 7 min read on the page · deepstack.ltd/stories/anthropic-grew-142-percent-openai/

THE DEEPSTACK CALL

Split Split the trade

Stop pricing OpenAI and Anthropic as one bet. One revenue line holds under pressure; the other has to be bought.

Conviction Moderate (3/5) · Version 2 · Oct 9, 2026

WHAT TO DO WITH THIS

1. Separate the two in any AI exposure: pricing power, not headline growth, sets the value.
2. Demand life-of-cluster margins in any prospectus, not first-year GPU economics.
3. Treat customer concentration and investor-funded revenue as the first red flags.

WHAT WOULD CHANGE OUR MIND

- Anthropic is forced into an OpenAI-style price cut on coding.
- Audited losses or a DeepSeek-style price reversal erase the gap between the two businesses.

THE FULL ARGUMENT

DeepStack's view: the market should stop treating OpenAI and Anthropic as a single economic object. The decisive evidence is not the headline growth gap but the way revenue behaves under pressure: OpenAI cut price 80% and lost $20.9bn, while Anthropic's enterprise coding share and $559m adjusted operating income suggest a narrower, more defensible product. That view remains conditional. Jeremie Eliahou Ontiveros's $100bn-per-GW-year calculation is powerful, but Daniel Romero's depreciation objection attacks the exact period that makes it look so good. A prospectus showing durable life-of-cluster margins, low customer concentration and no investor-funded revenue would strengthen the bulls. A comparable Anthropic coding cut, a DeepSeek reversal or audited losses would move us toward the bear case.

CALL HISTORY

Version	Date	Stance	Conviction	Reason	By
v2	Oct 9, 2026	Split the trade	Moderate (3/5)	Side named by the owner.	Marco Cauduro
v1	Sep 11, 2026	Split the trade	Moderate (3/5)	First call, as published with the story. The conviction, the actions and the flips were set down for the site on October 7, 2026.	DeepStack

BY THE NUMBERS

142%
Anthropic quarterly growth
Quarter on quarter, Q2 2026

18%
OpenAI quarterly growth
Quarter on quarter, Q2 2026

$20.9B
OpenAI 2025 loss
Steve Eisman; 2025

$559M
Anthropic adjusted income
Michael Spencer; preliminary June quarter 2026

WHO IS ON EACH SIDE

What does a frontier lab sell: a commodity, or a premium?

Each voice is attributed and paraphrased, dated where the date is known, with a link where one exists. DeepStack shows no score, ranking or accuracy for any person. The page: deepstack.ltd/sides/anthropic-grew-142-percent-openai/

4 VOICES

OpenAI: commoditized

Michael Spencer Aug 20, 2026

Michael Spencer predicts Anthropic will achieve AI supremacy status in 2027 and by 2030 will have no comparable peer competitor, not Google, OpenAI, DeepSeek, or Alibaba.

www.ai-supremacy.com/p/we-have-never-seen-an-anthropic-before-2026-revenue-ipo

Michael Spencer Aug 20, 2026

Open-source Chinese models are disrupting OpenAI's revenue because routing and code wrappers increasingly favor them over OpenAI's models due to cheaper token costs, forcing OpenAI to cut API prices, including an 80% cut to GPT-5.6 Lune, which…

www.ai-supremacy.com/p/we-have-never-seen-an-anthropic-before-2026-revenue-ipo

Massif Capital Jul 30, 2026

Chinese open-weight models climbed from under 2% of token traffic in late 2024 to a majority of usage among top models in 2026, driven by the revealed preference for models running on infrastructure the user controls.

fiscal.ai/fund-letters/massif-capital-llc-q2-2026-fscl_document_3AFW4FBM367

Massif Capital Jul 30, 2026

In June 2026, the Commerce Department suspended access to Anthropic's Fable 5 and Mythos 5 models for all non-US persons, invoking export-control authority, with practical effect of a global suspension that lasted under three weeks before…

fiscal.ai/fund-letters/massif-capital-llc-q2-2026-fscl_document_3AFW4FBM367

THE STRONGEST CASE AGAINST THIS SIDE

DeepValue Capital Aug 21, 2026

Alibaba's AI models have reached 3 billion downloads, surpassing Meta and Google.

Contradicts Michael Spencer on evidence · deepvaluecapitalbykyler.substack.com/p/week-ending-82126-news-sentiment

4 VOICES

Anthropic: not yet

Christian Catalini Aug 11, 2026

Training runs account for only 10 to 23 percent of AI labs' total compute costs, with the remainder tied to tacit knowledge in data curation, infrastructure, and learning from failed experiments.

www.a16z.news/p/some-simple-economics-of-open-versus

Christian Catalini Aug 11, 2026

The demand curve for AI capabilities is heterogeneous: commodity intelligence will become price-competitive across many applications, while convex-payoff domains like finance and frontier research can sustain premium pricing for marginal capability...

www.a16z.news/p/some-simple-economics-of-open-versus

Michael Spencer Aug 20, 2026

Anthropic's ARR grew from approximately $700 million in August 2024 to its current level, a roughly 93x increase in about two years.

www.ai-supremacy.com/p/we-have-never-seen-an-anthropic-before-2026-revenue-ipo

Michael Spencer Aug 20, 2026

Cognition's Devin achieved a $1 billion annualized revenue run rate and a $40 billion valuation in under two years, making it exceptional among AI startups for this pace of revenue growth.

www.ai-supremacy.com/p/visions-of-ai-automating-repetitive-grunt-coding-tasks-2026-devin

6 VOICES

Voices that cross sides

Michael Spencer Aug 20, 2026

As Devin and similar tools scale, they drive increased token usage of Anthropic's models, creating a compounding benefit to Anthropic.

www.ai-supremacy.com/p/visions-of-ai-automating-repetitive-grunt-coding-tasks-2026-devin

Michael Spencer Sep 10, 2026

Europe's dream of sovereign AI leadership is fading because Mistral, despite raising 3.5 billion dollars at a 21 billion euro valuation, has been unable to compete with the United States or China in open-source AI models.

www.ai-supremacy.com/p/ai-reality-european-sovereign-ai-is-dying-mistral-series-d

Wyndo Aug 13, 2026

Wyndo cites Artificial Analysis benchmark results as evidence of rapid open-weight model improvement: Kimi K3 scored 57 on its Intelligence Index (13 points above Kimi K2.6, per the July 17 evaluation), and DeepSeek V4 Flash 0731 scored 10 points...

aimaker.substack.com/p/claude-code-multiple-models-openrouter

Wyndo Aug 13, 2026

The traditional split of using a strongest frontier model (Opus) for reasoning/judgment and a faster cheaper model (Sonnet) for execution has become less clearly justified because the capability gap between 'thinking' and 'executing' models has…

aimaker.substack.com/p/claude-code-multiple-models-openrouter

Ed Elson Aug 11, 2026

Because OpenAI and Anthropic pay their bills with venture funding of which Big Tech supplies roughly 40%, revenue Big Tech books from them is partly recycled from itself and not real growth.

www.profgmedia.com/p/plastic-surgery

Ed Elson Aug 11, 2026

If Big Tech depends heavily on OpenAI and Anthropic for revenue growth, that growth is unsustainable because both customers are massively unprofitable and may not survive.

www.profgmedia.com/p/plastic-surgery

NOT A DISAGREEMENT

The number three bears collided with

SemiAnalysis: on a current cluster the labs can earn over $100B per gigawatt-year against about $12B of rental cost. Daniel Romero: a GPU earns most of its cash in year one while the books depreciate it straight. Steve Eisman: there is no dataset here like the subprime delinquency data; he waits for the two labs to open their books.

WHERE THEY CROSS

The specific points where one voice meets another: the same evidence read two ways, or the same mechanism with a different sign.

Ed Elson qualifies Michael Spencer on condition

Michael Spencer contradicts Ed Elson on premise

Michael Spencer qualifies Wyndo on condition

THE DATED TESTS

What would settle it, on the Docket. A test is specified when an outsider can settle it without asking anyone: the question, the source, the rule and what happens if the source is late, revised or silent.

Next prospectus p05-t01 PENDING

An audited listing reveals margins, customer concentration and investor-linked revenue; clean disclosure would support separation, weak economics would reunite the risks.

Specification pending

Next disclosures p05-t02 PENDING

Useful-life assumptions show whether first-year GPU margins survive across the cluster; Romero’s accounting would gain force if they do not.

Specification pending

DeepSeek repricing p05-t03 PENDING

A quiet reversal would weaken Erkan's demand evidence; sustained higher prices would support the view that demand remains structurally strong.

Specification pending

Jan 3, 2028 p05-t04 PENDING

Return on invested capital should be visible by Romero's deadline; failure would turn today's extraordinary economics into a first-year illusion.

Specification pending

A second revocation p05-t05 PENDING

Another model suspension would turn Massif's signal into procurement policy, showing that open weights are changing buying behavior.

Specification pending

THE TEST CAME IN

No test from this Part has resolved yet. Results are promoted by a person and listed on the Results page.

SOURCES AND FURTHER READING

1. AI Supremacy - Q2 lab revenue comparison · www.ai-supremacy.com/p/we-have-never-seen-an-anthropic-before-2026-revenue-ipo
2. Fiscal.ai - Massif Capital Q2 letter · fiscal.ai/fund-letters/massif-capital-llc-q2-2026-fscl_document_3AFW4FBM367
3. a16z News - Open versus closed AI economics · www.a16z.news/p/some-simple-economics-of-open-versus
4. Capitalist Letters - Frontier compute economics · www.capitalist-letters.com/p/its-time-to-be-bullish-on-compute
5. Hypertech Invest - AI selloff analysis · hypertechinvest.com/p/my-thoughts-on-the-ai-selloff-and
6. Financial Modeling Prep - Palantir Q2 call · financialmodelingprep.com/earning-call-transcript/PLTR/2026-q2

How DeepStack settles a question

Every Part of The Buildout asks one question and gathers the strongest case on each side from named investors and analysts, attributed, dated and linked. DeepStack then takes a side in The DeepStack Call, names what it would take to change its mind, and sets dated tests that an outsider can read on the day. When a test comes in, the result is published with its evidence, confirmed, refuted or inconclusive, and the Call's history records when the house changed its mind and why. No person is scored or ranked. No accuracy rate is shown until ten results exist.

Quoted passages are short, attributed by name and linked to their source; the paraphrases are DeepStack's. Quote this memo with attribution and a link to the canonical page.

This memo reflects the page as of October 12, 2026. Canonical page: https://deepstack.ltd/stories/anthropic-grew-142-percent-openai/. Who is on each side: https://deepstack.ltd/sides/anthropic-grew-142-percent-openai/.

Analysis and opinion, not investment advice.