

MARKETS

The Model Is Not the Price

A frontier model now lists at \$4 per million input tokens, open and Chinese models carry most of the tokens on the platforms that count them, and the two labs' revenue is counted in the tens of billions. Investors agree the token is cheap. They differ on what the lab is paid for, and only a prospectus on EDGAR can settle it.

By DeepStack · October 9, 2026 · Markets · 10 min read on the page · deepstack.ltd/stories/the-model-is-not-the-price/

THE DEEPSTACK CALL

Mixed The price has left the model

A frontier model no longer commands a price per token for long. The lab still does, and only a prospectus says what that price is worth.

Conviction Moderate (3/5) · Version 2 · Oct 10, 2026

WHAT TO DO WITH THIS

1. Separate token volume from token dollars:
OpenRouter's 57% to 67% and Vercel's 62% in September are volume; the dollar share is uncounted.
2. Date every run-rate before comparing two:
Longrivers's May, Romero's July and Kim's October 2 figures are not updates of one another.
3. Read the customer note before the growth rate when the prospectus lands: Silverio's quarter from two customers is the line to clear first.

WHAT WOULD CHANGE OUR MIND

- A prospectus with audited revenue near the relayed run-rates, no customer above a tenth of sales and no investor-linked revenue to disclose.
- Alphabet sizes the TPU contracts with Anthropic inside them on October 28 and Cloud growth holds without them, retiring Burry's charge.
- No registration on EDGAR by October 31 while either lab cuts list prices again; the house moves toward the open-weights reading.

THE FULL ARGUMENT

DeepStack's view is that the price has left the model. The gap to open weights is about four months by Slok's July 16 reading of Epoch, a majority of tokens on the two platforms that count them had moved to open or Chinese models by September, and Anthropic's list price undercut OpenAI's by 60% on September 27. What still commands a price is the lab. Ren measured it in the completed job, George located it in distribution, and Erkan counts it in the dollars that have not followed the volume. The single most important reason for caution is that no audited figure for that price is on EDGAR. Longrivers's \$50 billion in May, Romero's \$65 billion in July and Kim's \$70 billion for OpenAI on October 2 are three readings on three dates, and only the registration statement the Docket tests for by October 31 puts them on one line. DeepStack expects the filing to show a lab that is paid, in part by its investors: the first reading on the revenue, the third on what it is worth. The view changes if a prospectus shows audited revenue near the relayed run-rates with no customer above a tenth of sales and no investor-linked revenue large enough to

disclose, or if Alphabet sizes the TPU contracts on October 28 and Cloud growth holds without them. The first answer arrives October 28.

CALL HISTORY

Version	Date	Stance	Conviction	Reason	By
v2	Oct 10, 2026	The price has left the model	Moderate (3/5)	Side named by the owner.	Marco Cauduro
v1	Oct 9, 2026	The price has left the model	Moderate (3/5)	First call, with the story.	DeepStack

BY THE NUMBERS

\$4 vs \$10

Anthropic's Opus 5.5 list price per million input tokens, against OpenAI's GPT-6 Astra Charlie Hills, Sep 27, 2026; the frontier undercut by 60%

57% to 67%

Chinese models' share of token usage on OpenRouter by September As counted by Michael Spencer, Oct 1, 2026; Le Shrub had 66% on Sep 22

4 months

The gap between the Chinese open models and the frontier Torsten Slok of Apollo, citing the Epoch Capabilities Index, Jul 16, 2026

A quarter

Share of Anthropic's 2025 revenue that came from two customers Mia Silverio of Prof G Markets, Oct 5, 2026; 47% of sales ran through Amazon and Google

WHO IS ON EACH SIDE

Does the frontier model still command a price?

Each voice is attributed and paraphrased, dated where the date is known, with a link where one exists. DeepStack shows no score, ranking or accuracy for any person. The page: deepstack.ltd/sides/the-model-is-not-the-price/

11 VOICES

The frontier is still paid

Longriver Jul 20, 2026

Reported revenue acceleration at the two leading model companies — an annualised run-rate said to go from USD 1 billion at end-2024 to USD 9 billion at end-2025 to almost USD 50 billion in May this year at Anthropic, with a similar trajectory at...

fiscal.ai/fund-letters/longriver-investment-partners-limited-q2-2026-fscl_document_8HDF9YVG683

Tae Kim Oct 2, 2026

Tae Kim anticipates a close back-and-forth for model leadership between OpenAI and Anthropic over the next six months, expects another Anthropic model release before or around its IPO, and predicts both American frontier labs should do great once...

taekim.substack.com/p/openais-undeniable-resurgence-and

Tae Kim Oct 2, 2026

Tae Kim argues the earlier media consensus that OpenAI had fallen too far behind Anthropic was wrong, and that OpenAI pivoted resources to agentic coding after Anthropic's December lead, largely caught up by April, and arguably surpassed Anthropic...

taekim.substack.com/p/openais-undeniable-resurgence-and

Oguz Erkan Sep 13, 2026

Oguz Erkan says the concentration of the LLM market is still unknown: OpenAI and Anthropic take the lion's share of token dollars but their share of token volume has been declining, so the LLM market could be far more fragmented in 3-5 years, which...

www.capitalist-letters.com/p/booking-holdings-deep-dive-an-attractive

Wyndo Oct 4, 2026

Wyndo believes an always-on agent should run on the most capable model available, which is why they consider it fitting that ChatGPT's Dots runs on GPT-6 Astra.

Wyndo Sep 13, 2026

Wyndo recommends trying open-weight models if a user keeps hitting usage limits on Claude or ChatGPT, or wants to lower the cost of certain tasks, since some work may run well on an open-weight model while saving the subscription for tasks where...

aimaker.substack.com/p/claude-code-free-openrouter

Ren Sep 15, 2026

What actually fell is cost per completed job rather than cost per token: against Sol, Astra finishes Terminal-Bench 4.0 tasks at roughly 9% lower estimated API cost per task and OSWorld 2.0 tasks in about 47% less time, so the relevant demand...

renstocks.substack.com/p/openai-astra-has-agi-arrived

David George Sep 28, 2026

They judge the switching-cost lever basically unavailable because the AI frontier is a Red Queen's Race at the model layer, with swapping between models easy down to task-by-task switching and everyone improving continuously.

www.a16z.news/p/openai-understands-something-important

David George Sep 28, 2026

They say a pretty small group uses AI rampantly for everything and a somewhat larger group uses ChatGPT regularly, meaning most people have not yet been awakened as customers, and they think OpenAI understands how to do it.

www.a16z.news/p/openai-understands-something-important

Gaetano Sep 23, 2026

Criteo launched an Agentic Commerce Recommendation Service letting AI systems use Criteo's commerce intelligence to rank products, and OpenAI selected Criteo as its first advertising technology partner for ChatGPT Ads, with more than 2,000 brands...

Market Sentiment Sep 17, 2026

Based on their capex sustainability model, Market Sentiment estimates that next year's ~\$1.2T+ of AI capex should drive ~\$2T of incremental revenue, which they say will only be a fraction of the total incremental revenue AI generates as a business...

www.marketsentiment.co/p/the-next-great-ai-trade-is-everything

THE STRONGEST CASE AGAINST THIS SIDE**Michael Spencer** Oct 1, 2026

Anthropic is far ahead of OpenAI pre-IPO, both in model quality and ARR, and that lead should only accelerate with IPO proceeds heading into its November IPO.

Contradicts Tae Kim on evidence · www.ai-supremacy.com/p/ai-frontier-google-gemini-4-argon

10 VOICES**Open weights set the price****Michael Spencer** Oct 1, 2026

Chinese AI models captured a majority share of 57% to 67% of total token usage on the OpenRouter developer platform by September 2026, eroding OpenAI's token marketshare even as OpenAI subsidizes its new cheaper GPT-6.1 Sol model.

www.ai-supremacy.com/p/ai-frontier-google-gemini-4-argon

Michael Spencer Oct 1, 2026

Anthropic is far ahead of OpenAI pre-IPO, both in model quality and ARR, and that lead should only accelerate with IPO proceeds heading into its November IPO.

www.ai-supremacy.com/p/ai-frontier-google-gemini-4-argon

Steve Eisman Sep 11, 2026

OpenAI's price cut on its GPT model, combined with a reported tenfold jump in usage, looks like the beginning of a price war in AI services.

realeismanplaybook.substack.com/p/ai-terminator-fears-grow-and-rates

Steve Eisman Sep 28, 2026

Steve says the AI story probably continues as long as token growth keeps rising, and identifies a price war between Anthropic, OpenAI and the Chinese models as the main risk that could change it.

realeismanplaybook.substack.com/p/what-the-ai-boom-looks-like-from

Le Shrub Sep 22, 2026

China is gaining ground in AI, with Chinese models now accounting for 66% of token share.

www.shrubstack.com/p/uncomfortable-truths

Massif Capital

Under 2% of tokens to a majority.

Kevin Gee Sep 12, 2026

China is releasing open frontier models not despite state involvement, but because an open frontier is how you seed an ecosystem, win the developer base, and move the world's default tech stack onto a Chinese one, motivated by technological...

Deborah Folloni

100,000 leads: \$7 against \$3,000.

Charlie Hills Sep 27, 2026

Opus 5.5 is 60% cheaper than Astra, at \$4 versus \$10 per million input tokens.

charliehills.substack.com/p/claude-code-motion-graphics

David George Sep 30, 2026

So far, neither compute nor model advancements have been a zero-sum game: better, cheaper intelligence is accruing value to the entire ecosystem, with both older chips and models retaining substantial value well past the expiration dates the bears...

www.a16z.news/p/state-of-markets-ii

11 VOICES

The filing sets the price

Mia Silverio Oct 5, 2026

Anthropic's expected ~\$2 trillion IPO valuation is not justified by its numbers, given that it would need \$1.2 trillion in annual revenue within 10 years (per Aswath Damodaran) to justify that valuation while the entire current AI products and...

www.profgmedia.com/p/americans-are-turning-on-ai-anthropic

Mia Silverio Oct 5, 2026

Anthropic's customer concentration is a risk for its \$2 trillion valuation: nearly a quarter of its 2025 revenue came from two customers and 47% of sales ran through Amazon and Google, which is an acceptable problem for a small company but not for...

www.profgmedia.com/p/americans-are-turning-on-ai-anthropic

Distillate Capital Partners LLC Jul 15, 2026

Circular financing distorts the picture: hyperscalers and NVIDIA have committed roughly \$180 billion to OpenAI and Anthropic, which pledge about \$1.1 trillion of compute spending back to them against far smaller revenues, and \$53 billion of...

fiscal.ai/fund-letters/distillate-capital-partners-llc-q2-2026-fscl_document_3BZX2NEX647

Michael Burry Sep 24, 2026

Alphabet's reported 112% Google Cloud backlog growth is suspect because a new category of unknown size — TPU system sales contracts — was added to cloud backlog in the last six months, and one of the largest TPU customers is Anthropic, in which...

michaeljburry.substack.com/p/capital-cycle-iq-and-the-forensic

Michael Burry Sep 26, 2026

Aggregate purchase commitments, future leases, guarantees backing third-party debt, and various off-balance sheet potential liabilities across the hyperscalers total about \$3 trillion, pointing toward a couple more years of record AI data center...

Warden Capital Jul 22, 2026

Warden Capital argues the AI investment cycle cannot earn a reasonable return absent AGI within a few years: at roughly \$700 billion per year of investment, model providers would need about \$1 trillion in annual end AI revenue by 2029, ramping...

fiscal.ai/fund-letters/warden-capital-q2-2026-fscl_document_8ABS2CYK873

First Pacific Advisors Jul 30, 2026

Global AI spending is projected to reach 2.59 trillion dollars in 2026, and to achieve even a relatively low unlevered return on investment would require 207 billion dollars in after-tax income per year starting immediately, which the portfolio...

fiscal.ai/fund-letters/first-pacific-advisors-lp-q2-2026-fscl_document_9KDV3DNH810

Daniel Romero Sep 21, 2026

Although SpaceX reports AI infrastructure payback periods of less than one year, maintaining a one-year payback across the entire hypothetical 2027 buildout would require \$1.9-2.28 trillion in annual compute payments, and if AI labs footed that...

hypertechinvest.com/p/the-math-behind-the-next-ai-rally

Daniel Romero Sep 21, 2026

OpenAI's combined annualized revenue with Anthropic reached \$105 billion (OpenAI \$40 billion, Anthropic \$65 billion) in July 2026, which is nowhere near the figure needed to support the modeled 2027 buildout, and this understates the true gap since...

hypertechinvest.com/p/the-math-behind-the-next-ai-rally

Steve Eisman Sep 21, 2026

Anthropic looks likely to file for an IPO soon, and it will want to do that while token usage numbers are still near their peak, before the picture gets less flattering.

realeismanplaybook.substack.com/p/stock-picking-in-difficult-times

Steve Eisman Sep 21, 2026

The AI trade's foundations are shaky because roughly seventy percent of the hyperscalers' AI revenue traces back to just two customers, OpenAI and Anthropic, both of which are burning enormous amounts of cash, so if either stumbles the entire...

realeismanplaybook.substack.com/p/stock-picking-in-difficult-times

THE STRONGEST CASE AGAINST THIS SIDE**Market Sentiment** Sep 17, 2026

Based on their capex sustainability model, Market Sentiment estimates that next year's ~\$1.2T+ of AI capex should drive ~\$2T of incremental revenue, which they say will only be a fraction of the total incremental revenue AI generates as a business...

Contradicts Michael Burry on evidence · www.marketsentiment.co/p/the-next-great-ai-trade-is-everything

7 VOICES

Voices that cross sides

Michael Spencer Sep 29, 2026

OpenAI is in deep trouble in 2027 due to pressure in consumer AI from Meta and SpaceXAI, pressure from Chinese open-source models on cost and token marketshare, and being late to IPO while being overtaken in ARR and growth by Anthropic in...

www.ai-supremacy.com/p/the-era-of-personal-super-intelligent-agents-muse-2026

Michael Spencer Sep 11, 2026

Michael Spencer argues that in 2026 Chinese open-weight models are taking serious market share from incumbents like OpenAI and Anthropic, because the Chinese labs are only a few months behind the frontier and shortening release gaps mean even...

www.ai-supremacy.com/p/ai-insights-is-ai-adoption-slowng-down-ai-risk-debate

David George

Cheaper intelligence lifts all; at the model layer no edge sticks.

www.a16z.news/p/openai-understands-something-important

Steve Eisman Sep 11, 2026

Oracle's remaining performance obligations of \$664 billion are concerning because roughly 50% come from OpenAI, a company burning massive cash, making Oracle over-reliant on OpenAI's solvency.

realeismanplaybook.substack.com/p/ai-terminator-fears-grow-and-rates

Tae Kim

\$70B of ARR at OpenAI; Huang's 80% of AI-native startups on open models.

taekim.substack.com/p/openais-undeniable-resurgence-and

Le Shrub

Chinese models at 66% of tokens; Karp's never-IPO call on OpenAI.

www.shrubstack.com/p/uncomfortable-truths

Oguz Erkan

The dollars with two labs, the volume leaving, fragmentation in 3-5 years.

www.capitalist-letters.com/p/booking-holdings-deep-dive-an-attractive

NOT A DISAGREEMENT

The thread

Part 5 split the two labs into two businesses. Part 12 asks whether the model they sell still commands a price, and finds the writers agreeing on the volume leaving the frontier and splitting on the dollars.

WHERE THEY CROSS

The specific points where one voice meets another: the same evidence read two ways, or the same mechanism with a different sign.

Michael Spencer contradicts Tae Kim on evidence

Michael Spencer contradicts Le Shrub on evidence

Michael Spencer qualifies Wyndo on condition

Michael Spencer qualifies Oguz Erkan on evidence

Michael Spencer contradicts Le Shrub on premise

Michael Spencer qualifies Tae Kim on evidence

David George qualifies Michael Spencer on mechanism

Michael Spencer qualifies Tae Kim on condition

David George qualifies Steve Eisman on mechanism

Michael Spencer contradicts Wyndo on evidence

Michael Spencer contradicts David George on mechanism

Kevin Gee contradicts David George on mechanism

Tae Kim contradicts Michael Spencer on premise

Michael Spencer qualifies David George on mechanism

Gaetano qualifies Oguz Erkan on condition

David George qualifies Michael Spencer on condition

Daniel Romero contradicts Market Sentiment on evidence

Longriver contradicts Warden Capital on evidence

Longriver contradicts First Pacific Advisors on evidence

Longriver contradicts Distillate Capital Partners LLC on evidence

Michael Burry qualifies Daniel Romero on evidence

Michael Burry contradicts Market Sentiment on evidence

Gaetano qualifies Mia Silverio on condition

THE DATED TESTS

What would settle it, on the Docket. A test is specified when an outsider can settle it without asking anyone: the question, the source, the rule and what happens if the source is late, revised or silent.

Oct 28, 2026 p12-t01 PENDING

Alphabet, Microsoft and Meta report, and the Fed decides the same day: whether the TPU system sales inside Google Cloud's backlog get a size, and whether Cloud growth holds near the second quarter's 82% without them.

Question	Did Alphabet, in its third-quarter 2026 results of October 28, 2026 (the earnings release, the 10-Q or the call), state a size for the TPU system sales contracts inside Google Cloud's backlog?
Source	Alphabet Q3 2026 earnings release, call transcript and Form 10-Q (SEC EDGAR, CIK 1652044) www.sec.gov/cgi-bin/browse-edgar?action=getcompany&CIK=0001652044&type=10-Q
Rule	Confirmed if the release, the 10-Q or the call transcript gives a dollar amount, or a share of the Cloud backlog, for TPU system sales contracts or for the Anthropic contracts; refuted if the quarter is reported with no such figure. Google Cloud's growth rate is read alongside (the Part watches whether it holds near the second quarter's 82% without those contracts) but does not decide the test.
If the source is late, revised or silent	A size given only as a range wider than two to one, or given for a period other than the backlog at quarter end, is inconclusive, reviewed by the owner. A report delayed past November 15, 2026 is inconclusive.

Oct 29, 2026 p12-t02 PENDING

Amazon reports: it does not break Anthropic out, so the test is what AWS discloses about its largest customers, against Silverio's October 5 count of 47% of Anthropic's sales through Amazon and Google.

Question	Did Amazon, in its third-quarter 2026 results of October 29, 2026 (the earnings release, the 10-Q or the call), disclose a customer or a group of customers at 10% or more of AWS revenue, or quantify Anthropic's share of AWS sales?
Source	Amazon Q3 2026 earnings release, call transcript and Form 10-Q (SEC EDGAR, CIK 1018724) www.sec.gov/cgi-bin/browse-edgar?action=getcompany&CIK=0001018724&type=10-Q
Rule	Confirmed if the concentration note of the 10-Q, the release or the call names a customer or a group of customers at 10% or more of AWS revenue, or states Anthropic's share of AWS sales; refuted if the quarter is reported with no such disclosure. The figure is read against Silverio's October 5 count of 47% of Anthropic's sales running through Amazon and Google.
If the source is late, revised or silent	A concentration disclosure given only for Amazon as a whole, without an AWS figure, is inconclusive, reviewed by the owner. A report delayed past November 15, 2026 is inconclusive.

Oct 31, 2026 p12-t03 PENDING

Anthropic's listing, Part 10's test as the Docket states it: a registration statement with audited revenue, net loss and customer concentration public on EDGAR, or shares trading; a confidential draft without an EDGAR filing is inconclusive.

Specification pending

Nov 18, 2026 p12-t04 **PENDING**

Nvidia reports: whether it names the labs among its largest customers, against Eisman's September 21 count of roughly 70% of the hyperscalers' AI revenue from the two.

Question	Did Nvidia, in its third-quarter fiscal 2027 results of November 18, 2026 (the earnings release, the 10-Q or the call), name OpenAI or Anthropic as a customer at 10% or more of revenue, or state the two labs' share of its revenue?
Source	Nvidia Q3 FY2027 earnings release, call transcript and Form 10-Q (SEC EDGAR, CIK 1045810) www.sec.gov/cgi-bin/browse-edgar?action=getcompany&CIK=0001045810&type=10-Q
Rule	Confirmed if the concentration note of the 10-Q, the release or the call names either lab as a direct or indirect customer at 10% or more of revenue, or states a share of revenue for the two labs together; refuted if the customers at 10% or more are described without naming either lab and no share for the labs is given. The figure is read against Eisman's September 21 count of roughly 70% of the hyperscalers' AI revenue coming from the two labs.
If the source is late, revised or silent	A disclosure that names a lab only as an indirect end customer without a share is inconclusive, reviewed by the owner. A report delayed past December 15, 2026 is inconclusive.

When it files p12-t05 **PENDING**

The prospectus: the gross margin on inference, customer concentration against Silverio's October 5 count of a quarter from two customers, and revenue tied to investors against Distillate's July 15 count of \$180 billion committed and \$1.1 trillion pledged back.

Specification pending

THE TEST CAME IN

No test from this Part has resolved yet. Results are promoted by a person and listed on the Results page.

SOURCES AND FURTHER READING

1. AI Supremacy - Gemini-4 Argon and the OpenRouter token count · www.ai-supremacy.com/p/ai-frontier-google-gemini-4-argon
2. Longriver Investment Partners - Q2 2026 letter · fiscal.ai/fund-letters/longriver-investment-partners-limited-q2-2026-fscl_document_8HDF9YVG683
3. Prof G Markets - Anthropic wants \$2 trillion anyway · www.profgmedia.com/p/americans-are-turning-on-ai-anthropic
4. Distillate Capital Partners - Q2 2026 letter · fiscal.ai/fund-letters/distillate-capital-partners-llc-q2-2026-fscl_document_3BZX2NEX647
5. Capitalist Letters - Booking Holdings deep dive, the token dollars · www.capitalist-letters.com/p/booking-holdings-deep-dive-an-attractive
6. a16z - OpenAI understands something important and rare · www.a16z.news/p/openai-understands-something-important
7. AI Supremacy - The era of personal super-intelligent agents · www.ai-supremacy.com/p/the-era-of-personal-super-intelligent-agents-muse-2026
8. Apollo - Chinese models vs. frontier models · www.apollo.com/insights-news/insights/daily-spark/chinese-models-vs-frontier-models
9. Shrubstack - Uncomfortable truths · www.shrubstack.com/p/uncomfortable-truths
10. Potential Multibaggers - Vercel's open-model count · www.potentialmultibaggers.com/p/and-another-new-pick
11. a16z - State of Markets II · www.a16z.news/p/state-of-markets-ii
12. First Pacific Advisors - Q2 2026 letter · fiscal.ai/fund-letters/first-pacific-advisors-lp-q2-2026-fscl_document_9KDV3DNH810
13. Warden Capital - Q2 2026 letter · fiscal.ai/fund-letters/warden-capital-q2-2026-fscl_document_8ABS2CYK873

14. Market Sentiment - The next great AI trade is everything that isn't AI · www.marketsentiment.co/p/the-next-great-ai-trade-is-everything
 15. HyperTech Invest - The math behind the next AI rally · hypertechinvest.com/p/the-math-behind-the-next-ai-rally
 16. AI Supremacy - A guide to Grok Bot 2026, the listing delay · www.ai-supremacy.com/p/a-guide-to-grok-bot-2026-autonomous-teammate-systems
-

How DeepStack settles a question

Every Part of The Buildout asks one question and gathers the strongest case on each side from named investors and analysts, attributed, dated and linked. DeepStack then takes a side in The DeepStack Call, names what it would take to change its mind, and sets dated tests that an outsider can read on the day. When a test comes in, the result is published with its evidence, confirmed, refuted or inconclusive, and the Call's history records when the house changed its mind and why. No person is scored or ranked. No accuracy rate is shown until ten results exist.

Quoted passages are short, attributed by name and linked to their source; the paraphrases are DeepStack's. Quote this memo with attribution and a link to the canonical page.

This memo reflects the page as of October 10, 2026. Canonical page: <https://deepstack.ltd/stories/the-model-is-not-the-price/>. Who is on each side: <https://deepstack.ltd/sides/the-model-is-not-the-price/>.

Analysis and opinion, not investment advice.